

# ACID and Indexes:

In MongoDB, ACID (Atomicity, Consistency, Isolation, Durability) transactions are supported to ensure that database operations are executed reliably.

Regarding indexes, MongoDB supports creating collections and indexes in transactions implicitly or explicitly.

## Atomicity

Atomicity ensures that database transactions are treated as a single, indivisible unit of work. If any part of the transaction fails, the entire transaction is rolled back, and the database is returned to its previous state.

### Example in MongoDB:

Suppose we have a banking application that transfers money from one account to another. We want to ensure that either both accounts are updated successfully or neither is updated if an error occurs.

```
// Start a session
```

```
session = client.startSession();
```

```
// Start a transaction
```

```
session.startTransaction();
```

```
try {
```

```
    // Debit $100 from account A
```

```
    db.accounts.updateOne({ _id: "A" }, { $inc: { balance: -100 } });
```

```
    // Credit $100 to account B
```

```
    db.accounts.updateOne({ _id: "B" }, { $inc: { balance: 100 } });
```

```
// Commit the transaction
session.commitTransaction();
} catch (error) {
    // Abort the transaction if an error occurs
    session.abortTransaction();
}
```

In this example, if the debit operation fails, the credit operation will not be executed, and the transaction will be rolled back.

## Consistency

Consistency ensures that the database remains in a valid state, even after multiple transactions have been executed. This means that the database must always satisfy certain constraints, such as data types, relationships, and business rules.

### Example in MongoDB:

Suppose we have a collection of orders with a total value that must always be greater than or equal to zero.

```
// Create an order with a total value of $100
db.orders.insertOne({ _id: 1, total: 100 });

// Update the order to have a total value of -$50
try {
    db.orders.updateOne({ _id: 1 }, { $set: { total: -50 } });
} catch (error) {
    // The update will fail because the total value cannot be negative
}
```

```
console.log(error);  
  
}
```

In this example, the update operation will fail because it would violate the consistency constraint that the total value must be greater than or equal to zero.

## Isolation

Isolation ensures that concurrent transactions do not interfere with each other. Each transaction sees a consistent view of the database, as if it were the only transaction executing.

### Example in MongoDB:

Suppose we have two concurrent transactions that update the same document.

```
// Transaction 1: Update the document to have a value of 10
```

```
session1.startTransaction();
```

```
db.collection.updateOne({ _id: 1 }, { $set: { value: 10 } });
```

```
session1.commitTransaction();
```

```
// Transaction 2: Update the document to have a value of 20
```

```
session2.startTransaction();
```

```
db.collection.updateOne({ _id: 1 }, { $set: { value: 20 } });
```

```
session2.commitTransaction();
```

In this example, each transaction sees a consistent view of the document, as if it were the only transaction executing. The final value of the document will be 20, because Transaction 2 committed last.

## Durability

Durability ensures that once a transaction has been committed, its effects are permanent and cannot be rolled back. This means that the database must guarantee that the data is safely stored and will not be lost in case of a failure.

### **Example in MongoDB:**

Suppose we have a transaction that updates a document and then crashes before the update is written to disk.

```
// Start a transaction

session.startTransaction();

try {

    // Update the document

    db.collection.updateOne({ _id: 1 }, { $set: { value: 10 } });

    // Commit the transaction

    session.commitTransaction();

    // Crash before the update is written to disk

    process.exit(1);

} catch (error) {

    // Abort the transaction if an error occurs

    session.abortTransaction();

}
```

In this example, even though the transaction crashed before the update was written to disk, the update will still be persisted when the database is restarted. This is because MongoDB uses a write-ahead log to ensure durability.

## REPLICATION and SHARDING:

Replication and Sharding are two important features in MongoDB for scalability and data availability.

### Replication

Replication increases redundancy and data availability by maintaining multiple copies of data on different database servers. This increases the performance of read scaling. A set of servers that maintain the same copy of data is known as replica servers or MongoDB instances.

A replica set is a cluster of N different nodes that maintain the same copy of the data set. The primary server receives all write operations and records all the changes to the data, i.e., oplog. The secondary members then copy and apply these changes in an asynchronous process. All the secondary nodes are connected with the primary nodes. There is one heartbeat signal from the primary nodes. If the primary server goes down, an eligible secondary will hold the new primary.

#### Replication provides:

- High Availability of data disasters recovery
- No downtime for maintenance (like backups, index rebuilds, and compaction)
- Read Scaling (Extra copies to read from)

### Sharding

Sharding is a method for distributing large collections (datasets) and allocating it across multiple servers. It helps in horizontal scaling by partitioning the large collection into smaller discrete parts called shards.

Sharding is essential for handling large amounts of data and high traffic. It allows MongoDB to scale horizontally, which means we can add more servers to handle increased traffic and data.

By combining replication and sharding, we can achieve high availability, scalability, and performance in our MongoDB database.

## Replication vs Sharding

	Replication	Sharding
Purpose	Increase data availability and redundancy	Distribute large datasets across multiple servers
Goal	Provide high availability and fault tolerance	Scale horizontally to handle large amounts of data and high traffic
Method	Maintain multiple copies of data on different servers	Partition large collections into smaller, discrete parts (shards)
Data Distribution	All data is replicated to all nodes	Data is split across multiple nodes, each containing a portion of the data
Read Performance	Improved read performance due to multiple copies of data	Improved read performance due to data being distributed across multiple nodes
Write Performance	Write performance is not improved, as all writes go to the primary node	Write performance is improved, as writes can be distributed across multiple nodes
Scalability	Limited scalability, as all nodes must contain a full copy of the data	Highly scalable, as new nodes can be added to handle increased data and traffic
Complexity	Relatively simple to implement and manage	More complex to implement and manage, especially for large datasets
Use Cases	Suitable for small to medium-sized datasets, high availability, and disaster recovery	Suitable for large datasets, high traffic, and horizontal scaling

### Key differences:

- Purpose:** Replication is primarily used for high availability and fault tolerance, while sharding is used for horizontal scaling and handling large datasets.

2. **Data Distribution:** Replication involves maintaining multiple copies of the entire dataset on different nodes, while sharding involves partitioning the dataset into smaller parts and distributing them across multiple nodes.
3. **Write Performance:** Replication does not improve write performance, as all writes go to the primary node, while sharding can improve write performance by distributing writes across multiple nodes.
4. **Scalability:** Replication has limited scalability, as all nodes must contain a full copy of the data, while sharding is highly scalable, as new nodes can be added to handle increased data and traffic.

#### When to use each:

- Use replication when:
  - You need high availability and fault tolerance for your dataset.
  - Your dataset is relatively small to medium-sized.
  - You want to ensure disaster recovery and business continuity.
- Use sharding when:
  - You have a large dataset that needs to be distributed across multiple nodes.
  - You need to handle high traffic and scale horizontally.
  - You want to improve write performance and reduce the load on individual nodes.

## Combining Replication and Sharding in MongoDB: (REPLICATION + SHARDING)

Replication and sharding are two powerful features in MongoDB that can be combined to achieve high availability, scalability, and performance. By combining replication and sharding, you can create a highly available and scalable MongoDB cluster that can handle large amounts of data and high traffic.

#### Benefits of Combining Replication and Sharding:

5. **High Availability:** Replication ensures that data is available even in the event of node failures, while sharding allows you to distribute data across multiple nodes, reducing the risk of data loss.
6. **Scalability:** Sharding allows you to scale horizontally by adding more nodes to handle increased data and traffic, while replication ensures that data is available on multiple nodes.

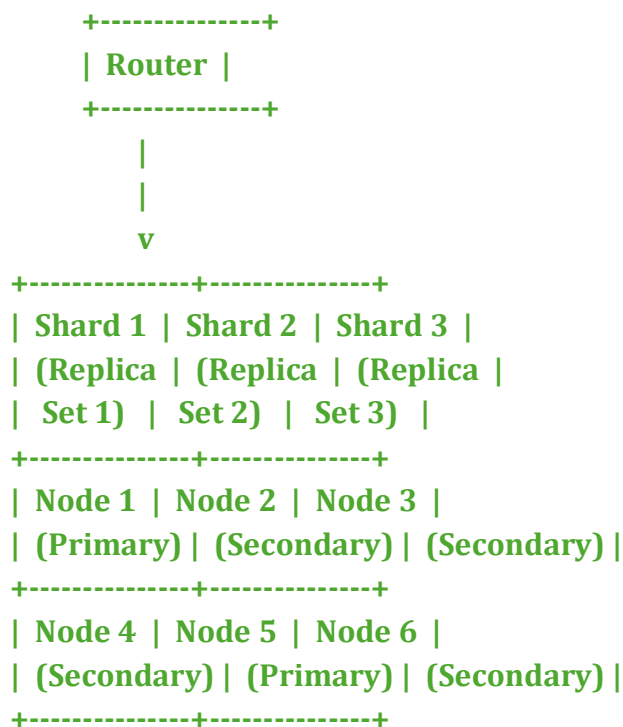
7. **Improved Performance:** By distributing data across multiple nodes, sharding can improve read and write performance, while replication ensures that data is available on multiple nodes, reducing the load on individual nodes.
8. **Disaster Recovery:** Combining replication and sharding provides a robust disaster recovery strategy, ensuring that data is available even in the event of node failures or data center outages.

### How to Combine Replication and Sharding:

9. **Create a Replica Set:** Create a replica set with multiple nodes, each containing a copy of the data.
10. **Shard the Data:** Shard the data across multiple nodes, using a shard key to determine which node to store the data on.
11. **Configure Sharding:** Configure sharding to distribute data across multiple nodes, using a shard key to determine which node to store the data on.
12. **Add Replication to Each Shard:** Add replication to each shard, ensuring that each shard has multiple nodes with a copy of the data.

### Example Architecture:

Here's an example architecture that combines replication and sharding:





In this example, we have three shards, each with a replica set containing three nodes. The router directs traffic to the appropriate shard, and each shard has multiple nodes with a copy of the data.

## INDEXES:

Indexes are a crucial component of MongoDB, allowing you to improve query performance, reduce latency, and optimize data retrieval. In this section, we'll explore the different types of indexes, how to create and manage them, and best practices for indexing in MongoDB.

### Types of Indexes in MongoDB:

- 13. **Single Field Index:** An index on a single field in a document.
- 14. **Compound Index:** An index on multiple fields in a document.
- 15. **Multi-Key Index:** An index on an array field, allowing for efficient querying of array elements.
- 16. **Text Index:** A specialized index for text search, allowing for efficient querying of text data.
- 17. **Hashed Index:** An index on a field that uses a hash function to map values to a fixed-size string.

### Special Indexes in MongoDB

In addition to the standard single-field and compound indexes, MongoDB provides several special indexes that cater to specific use cases and query patterns. These special indexes can significantly improve query performance and efficiency.

#### 1. Text Indexes

Text indexes are used for text search queries. They allow MongoDB to efficiently search for strings within a collection.

#### Creating a Text Index:

```
db.collection.createIndex({ field: "text" })
```

#### 2. Hashed Indexes

Hashed indexes are used for equality queries on fields with high cardinality (i.e., many unique values). They are particularly useful for fields with a large number of unique values.

#### **Creating a Hashed Index:**

```
db.collection.createIndex({ field: "hashed" })
```

### **3. TTL (Time-To-Live) Indexes**

TTL indexes are used to automatically remove documents from a collection after a specified time period. They are useful for implementing document expiration and caching.

#### **Creating a TTL Index:**

```
db.collection.createIndex({ field: 1 }, { expireAfterSeconds: 3600 })
```